

Job Req #1: Agentic AI Systems Engineer / Agentic AI Architect

We are looking for an Agentic AI Systems Engineer to design and build advanced agentic AI systems for chip design automation and engineering workflows. The role focuses on multi-agent orchestration, tool-calling systems, memory architectures, retrieval pipelines, long-horizon reasoning systems, and evaluation infrastructure. The ideal candidate should have 2–8+ years of experience in Applied AI, LLM Systems, AI Infrastructure, or ML Systems Engineering, along with strong Python and backend engineering skills.

Candidates must have hands-on experience building agentic AI systems involving planning, tool use, memory management, retrieval, autonomous execution, task decomposition, and feedback-driven workflows. Strong familiarity with frameworks such as LangChain, LangGraph, LlamaIndex, AutoGen, CrewAI, DSPy, Semantic Kernel, FastAPI, and cloud LLM APIs is required (or similar). The engineer should also be capable of building scalable evaluation infrastructure and dashboards for tracking metrics such as Pass@k, tool-call success rates, execution correctness, latency, hallucination/error rates, and regression metrics. Experience with LangSmith, Weights & Biases, MLflow, Grafana, Prometheus, Helicone, or similar observability and evaluation tooling is preferred.

Strong experience with RAG pipelines, vector databases, semantic retrieval, memory systems, and cloud-based AI infrastructure is highly desirable. Familiarity with Docker, Kubernetes, Redis, PostgreSQL, distributed task systems, and scalable backend infrastructure is preferred. This role requires a strong ownership mindset, fast execution, and the ability to rapidly iterate in a startup environment focused on advanced AI systems and autonomous engineering workflows.